

Release Captain

An evidence-led release workflow with a human approval gate.

Problem and workflow

Release teams reconcile commits, pull requests, diffs, tags and CI to choose versions and write notes. Scattered evidence makes signals easy to miss.

Users provide a repository, branch, optional tag and allowlisted test command. Live mode uses that tag, the latest tag or a default-branch comparison, then collects commits, linked pull requests, changed-file summaries and recent Actions runs. Commit-message rules categorize changes and flag breaking markers; removed files and CI failures become risks. Release Captain recommends a semantic version and drafts notes. The review shows evidence, risks, test status and audit events, then stops at **WAITING FOR HUMAN APPROVAL**. Approval must match the repository, version and action list.

Architecture and TrueForge

The Next.js dashboard uses FastAPI for GitHub collection, deterministic analysis, session state, approval validation, audit and execution. A checked-in TrueForge manifest and Python SDK connector can register or sync the agent and stream sessions when TrueForge and GitHub MCP are configured. Dashboard recommendations use deterministic FastAPI rules; its sessions do not pass through the connector.

What's implemented and current limits

Live GitHub analysis is implemented. After exact-plan approval, the backend can create a Git tag and GitHub release when execution is enabled and the token has write access. Demo evidence, passing tests and post-approval actions are simulated. TrueForge sandbox is disabled by default. The optional test runner executes allowlisted commands on local clones of public repositories, outside a sandbox. Version bumps follow commit prefixes and explicit breaking markers; removed files and CI failures appear as risks. These checks are limited; review each recommendation. Session and audit history are in memory; backend restart clears them. Publishing and deployment are not implemented.